

# dog\_app

November 18, 2019

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: \* Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dogImages.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

Reviewer comments 2: Please include the HTML report, you can get this by opening File > Download as... > HTML in the Jupyter notebook menu. PS. It's sufficient to only include the Jupyter notebook, HTML report file and the tested images from step 6.

Dear reviewer, I have published the jupyter notebook and HTML file at the git hub [dog breed](#) repository. Regarding the images used in step 6. I downloaded them and stored them in the repository `img` subdirectory. Git hub rejects publishing the original images sizes exceed the file size limit.

```
In [1]: import numpy as np
import pandas as pd
import cv2                                     # image processing
import torch
import torch.nn as nn
import torch.optim as optim                   # optimizer and loss function
import torch.nn.functional as F              # activation functions
import torchvision.models as models           # to load vgg16 model
import torchvision.transforms as transforms
import os                                     # system calls
import matplotlib.pyplot as plt              # plotting
import requests                               # for downloading data from a given URL
import ast
from tqdm import tqdm                         # show the progress of a loop, i.e tqdm
from glob import glob                         # image I/O
from PIL import Image                         # image manipulation
from PIL import ImageFile
from torchvision import datasets

import warnings
warnings.filterwarnings('ignore')

ImageFile.LOAD_TRUNCATED_IMAGES = True
%matplotlib inline

# evaluate if cuda is available
cuda_available = torch.cuda.is_available()

In [2]: # load filenames for human and dog images
human_files = np.array(glob("lfw/**/*.jpg"))
dog_files = np.array(glob("dogImages/**/*.jpg"))

# print number of images in each dataset
print('There are %d total human images.' % len(human_files))
print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.  
There are 8351 total dog images.

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [37]: # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)


         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))

         # get bounding box for each detected face
         for (x,y,w,h) in faces:
             # add bounding box to color image
             cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

         # convert BGR image to RGB for plotting
         cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

         # display the image, along with bounding box
         plt.imshow(cv_rgb)
         plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [28]: # returns "True" if face is detected in image stored at img_path
def Haar_detector(img_path):
    # Load the model
    face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?

- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

```
In [73]: # NOTE: Helper function
def detector_frequency(imgset, heuristic=None):
    bset = np.zeros((len(imgset)), dtype=bool)
    for i, fp in enumerate(imgset):
        bset[i] = heuristic(fp)
    return bset
```

**Answer:**

Reviewer comments 1: What percentage of the first 100 images in `human_files` have a detected human face?

96%

Reviewer comments 1 : What percentage of the first 100 images in `dog_files` have a detected human face?

18%.

**Short summary:**

100 images corresponding to the people and dogs were tested for human face detection. Results show, that 96% and 18% of people and dog images contained a human face. 4% of human faces were not detected and 18% of dogs were incorrectly classified as having a human face. Results might be explained since Haar classifier use change in contrast values between adjacent rectangular groups of pixels to detect features, such as eyes, mouth, and the edge of the face. However, many of these features are also present in dog and other classes of animal faces.

```
In [7]: human_files_short = human_files[:100]
dog_files_short = dog_files[:100]
```

```
In [75]: #-#-# Do NOT modify the code above this line. #-#-#
## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
hset = pd.DataFrame(detector_frequency(human_files_short, heuristic=Haar_detector))
dset = pd.DataFrame(detector_frequency(dog_files_short, heuristic=Haar_detector))
```

```
In [76]: print('Human : ', hset[hset[0]==True].shape[0]/hset.shape[0])
print('Dogs : ', dset[dset[0]==True].shape[0]/dset.shape[0])
```

Human : 0.96

Dogs : 0.18

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make

use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)  
       ### TODO: Test performance of another face detection algorithm.  
       ### Feel free to use as many code cells as needed.
```

---

## ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [40]: # define VGG16 model  
        vgg16 = models.vgg16(pretrained=True)  
  
        # check if CUDA is available  
        if cuda_available:  
            vgg16 = vgg16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [3]: # NOTE: Helper function  
        def vgg16_labels():  
            LABELS_MAP_URL = "https://gist.githubusercontent.com/yrevar/942d3a0ac09ec9e5eb3a/r  
            labels = ast.literal_eval(requests.get(LABELS_MAP_URL).text)  
            return labels
```

```
In [3]: def VGG16_predict(img_path):  
        '''  
        Use pre-trained VGG-16 model to obtain index corresponding to  
        predicted ImageNet class for image at specified path
```

```

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
    '''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    # load the model

    # define VGG16 model
    vgg16 = models.vgg16(pretrained=True)
    # check if CUDA is available
    if cuda_available:
        vgg16 = vgg16.cuda()

    img = Image.open(img_path)    # load the image from persistence

    # define image transformations
    transformations = transforms.Compose([
        transforms.Resize(size=224),
        transforms.CenterCrop((224,224)),
        transforms.ToTensor(),
        transforms.Normalize( mean=[0.485, 0.456, 0.406],
                               std=[0.229, 0.224, 0.225]))

    # apply the image transformations
    img = transformations(img).unsqueeze_(0)

    # map to gpu
    if cuda_available:
        img = img.cuda()

    # predict the class for the image
    output = vgg16(img)

    # convert output probabilities to predicted class
    _, pred = torch.max(output, 1)

    pred = np.squeeze(pred.numpy()) if not cuda_available else np.squeeze(pred.cpu().numpy())
    return int(pred)

```

In [8]: # Test code

```

labels = vgg16_labels()
cid = VGG16_predict(dog_files_short[10])
print(labels[cid])

```

Afghan hound, Afghan

```

In [48]: cid = VGG16_predict(human_files[3])
print(labels[cid])

```

neck brace

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog\_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [9]: ### returns "True" if a dog is detected in the image stored at img_path
def VGG16_dog_detector(img_path):
    ## TODO: Complete the function.
    cid = VGG16_predict(img_path)
    if(cid >=151 and cid<=268):
        return True
    return False

```

### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your dog\_detector function.

- What percentage of the images in human\_files\_short have a detected dog?
- What percentage of the images in dog\_files\_short have a detected dog?

**Answer:**

Reviewer comments 1: What percentage of the images in human\_files\_short have a detected dog?

None.

Reviewer comments 1: What percentage of the images in dog\_files\_short have a detected dog?  
95%.

```

In [78]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
hset = pd.DataFrame(detector_frequency(human_files_short, heuristic=VGG16_dog_detector))
dset = pd.DataFrame(detector_frequency(dog_files_short, heuristic=VGG16_dog_detector))

In [79]: print('Human : ', hset[hset[0]==True].shape[0]/hset.shape[0])
print('Dogs : ', dset[dset[0]==True].shape[0]/dset.shape[0])

```



Human : 0.0  
Dogs : 0.95

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

---

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

---

Brittany	Welsh Springer Spaniel
----------	------------------------

---

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

---

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

---

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

---

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

---

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

Reviewer comments 2: Write three separate data loaders for the training, validation, and test datasets of dog images. These images should be pre-processed to be of the correct size. The preprocessing code is good, the only problem is that you also applied data augmentation on the validation and test set. The reason we don't want data augmentation for validation and testing is that we want the validation and test sets to represent real images as much as possible. If we augment the images they might not be realistic anymore. For the transformation for the validation and test set it's sufficient to only resize the image to 224x224 pixels.

Dear reviewer. I appreciate the acute observation. Unconsciously I reused the same data loader for validation and test phases. I have written the two new data loaders, namely: the validation and test loaders. These only apply image resize (224,224) and normalization. I ran new experiments using the transfer learning model that previously achieved test precision of 84%. Follow this Section ?? for new results.

In [36]: *### TODO: Write data loaders for training, validation, and test sets*  
*## Specify appropriate transforms, and batch\_sizes*

```
data_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.Resize(size=224),
    transforms.CenterCrop((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))

data_transform_test = transforms.Compose([
    transforms.Resize(size=(224,224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))

data_transform_validate = transforms.Compose([
    transforms.Resize(size=224),
    transforms.CenterCrop((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))

data_dir = 'dogImages/'
batch = 20
workers=0
```

```

trainset = datasets.ImageFolder(os.path.join(data_dir, 'train/'), transform=data_transf
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch, num_workers=work

testset = datasets.ImageFolder(os.path.join(data_dir, 'test/'), transform=data_transf
testloader = torch.utils.data.DataLoader(testset, batch_size=batch, num_workers=work

validset = datasets.ImageFolder(os.path.join(data_dir, 'valid/'), transform=data_transf
validloader = torch.utils.data.DataLoader(validset, batch_size=batch, num_workers=work

loaders_scratch = dict()
loaders_scratch['train'] = trainloader
loaders_scratch['test'] = testloader
loaders_scratch['valid'] = validloader

```

In [32]: *# Helper function:*

```

def display_workload(loader):
    dataiter = iter(loader)
    imgset, lbset = dataiter.next()
    # imgset = imgset.numpy()


    fig = plt.figure(figsize=(25,4))
    for idx in np.arange(20):
        ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
        img = imgset[idx]
        img = (img-torch.min(img))/(torch.max(img)-torch.min(img)) #apply min-max norm
        img = np.transpose(img,(1, 2, 0))
        plt.imshow(img)

```

In [11]: display\_workload(loaders\_scratch['train'])



In [33]: display\_workload(loaders\_scratch['test'])



```
In [34]: display_workload(loaders_scratch['valid'])
```



**Question 3:** Describe your chosen procedure for preprocessing the data. 1. How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? 2. Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:**

Reviewer comments 1: How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?

Images were randomly rotated, resized to 224 x 224 and used 3 channels, centered crop, and normalized. Images are relatively large, compared with MNIST images sizes. The reason for working with the proposed size, was to preserve, as many as possible, the original image properties and invariants. So that, during training these are generalized by the learning heuristic. Secondly, I also considered the amount of available resources, CPU and GPU memory buffer sizes, as a criterion for bounding the image size. I tested several image sizes.

Reviewer comments 1: Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

There are numerous augmentation techniques, i.e: geometric transformations, color space augmentations, kernel filters, mixing images, random erasing, feature space augmentation, adversarial training, generative adversarial networks, among others. See (Connor Shorten and Taghi M. Khoshgohar, 2019). Augmentation techniques are often applied as a pre-processing phase, where produced images are added to the original dataset. Kernel filters are commonly applied to blur or sharpen images. Convolutional neural networks include kernel filters in its encoding phase. Thus it automatically includes data augmentation. I did not apply data augmentation as a pre-processing phase.

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [42]: # define the CNN architecture
class Net(nn.Module):
    """ TODO: choose an architecture, and complete the class """
    def __init__(self):
        super(Net, self).__init__()

        #define the convolutional layer
        self.conv1 = nn.Conv2d(3,16,3,padding=1)
        self.conv2 = nn.Conv2d(16,32,3,padding=1)
        self.conv3 = nn.Conv2d(32,64,3,padding=1)
```

```

self.pool = nn.MaxPool2d(2, 2)

## Define layers of a CNN
self.fc1 = nn.Linear(64*28*28,1000)
self.fc2 = nn.Linear(1000,1000)
self.fc3 = nn.Linear(1000,500)
self.fc4 = nn.Linear(500,133)

self.batch_norm2 = nn.BatchNorm1d(num_features=1000)
self.batch_norm3 = nn.BatchNorm1d(num_features=500)
self.dropout = nn.Dropout(0.25)

def forward(self, x):
    ## Define forward behavior
    x = self.pool(F.relu(self.conv1(x)))
    x = self.dropout(x)
    x = self.pool(F.relu(self.conv2(x)))
    x = self.dropout(x)
    x = self.pool(F.relu(self.conv3(x)))

    x = x.view(x.size(0), -1)

    x = F.relu(self.fc1(x))
    x = self.dropout(x)
    x = F.relu(self.fc2(x))
    x = self.dropout(x)
    x = self.batch_norm2(x)
    x = F.relu(self.fc3(x))
    x = self.batch_norm3(x)
    x = self.fc4(x)

    return x

### You do NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if cuda_available:
    model_scratch.cuda()

```

In [43]: `print(model_scratch)`

```

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(fc1): Linear(in_features=50176, out_features=1000, bias=True)
(fc2): Linear(in_features=1000, out_features=1000, bias=True)
(fc3): Linear(in_features=1000, out_features=500, bias=True)
(fc4): Linear(in_features=500, out_features=133, bias=True)
(batch_norm2): BatchNorm1d(1000, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(batch_norm3): BatchNorm1d(500, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(dropout): Dropout(p=0.25, inplace=False)
)

```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:**

Reviewer comments 1: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

The CNN consists of three convolutional layers whose aim is to augment the data set by both reducing image noise and increasing the acuteness of images. After each convolutional phase, negative values are removed via the relu activation function. Since the size and number of images resulting after a convolutional phase may be large, pooling is applied in order to reduce the amount of information and size of the image. Pooling may introduce information loss, but may be reduced by data augmentation. Finally, dropout is introduced in order to reduce overfitting. The depth, number of layers, in the reduction phase is bounded by the amount of information loss tolerable for the problem instance. In the proposed solution, images were shrunk significantly until their size reached 28x28.

The neural network layer, consists of four layers. Two dropout layers were added in order to reduce overfitting. Batch normalization is added in order to accelerate training and reduce covariance shift. Two phase of normalization was applied.

The network was tested only a few times, but results were not promising. A test accuracy of 13% was achieved. The proposed feature extraction phase depth, which compromise the convolutional layers, is shallow when compared to VGG16 feature extraction phase. Perhaps, the performance of the proposed network might be improved by improving such phase.

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [44]: ### TODO: select loss function
         criterion_scratch = nn.CrossEntropyLoss()

         ### TODO: select optimizer
         optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)

```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [45]: # the following import is required for training to be robust to truncated images
ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    # record train and validation losses
    loss_batch = []
    loss_valid = []

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_

            # clear the gradients of all optimized variables
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)
            # backward pass: compute gradient of the loss with respect to model param
            loss.backward()
            # perform a single optimization step (parameter update)
            optimizer.step()
            # update training loss
            train_loss += loss.item()*data.size(0)

        #####
        # validate the model #
        #####
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):
            # move to GPU

```



```

        if use_cuda:
            data, target = data.cuda(), target.cuda()
            ## update the average validation loss

            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)
            # update average validation loss
            valid_loss += loss.item()*data.size(0)

        # calculate average losses
        train_loss = train_loss/len(loaders['train'].sampler)
        valid_loss = valid_loss/len(loaders['valid'].sampler)
        #record average losses
        loss_batch.append(train_loss)
        loss_valid.append(valid_loss)

        # print training/validation statistics
        print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
            epoch,
            train_loss,
            valid_loss
        ))

        ## TODO: save the model if validation loss has decreased
        # save model if validation loss has decreased
        if valid_loss <= valid_loss_min:
            print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
                valid_loss_min,
                valid_loss))
            torch.save(model.state_dict(),save_path)
            valid_loss_min = valid_loss
    return loss_batch, valid_loss, model

# train the model
loss_batch, valid_loss, model_scratch = train(100, loaders_scratch, model_scratch, optimizer_scratch,
        criterion_scratch, cuda_available, 'model_scratch2.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch2.pt'))

```

Epoch: 1            Training Loss: 4.983277            Validation Loss: 126.376283  
Validation loss decreased (inf --> 126.376283). Saving model ...  
Epoch: 2            Training Loss: 4.807230            Validation Loss: 4.897680  
Validation loss decreased (126.376283 --> 4.897680). Saving model ...



Epoch: 3	Training Loss: 4.669751	Validation Loss: 4.891207
Validation loss decreased (4.897680 --> 4.891207). Saving model ...		
Epoch: 4	Training Loss: 4.511940	Validation Loss: 4.522735
Validation loss decreased (4.891207 --> 4.522735). Saving model ...		
Epoch: 5	Training Loss: 4.395568	Validation Loss: 4.641390
Epoch: 6	Training Loss: 4.297917	Validation Loss: 4.675839
Epoch: 7	Training Loss: 4.209301	Validation Loss: 4.440095
Validation loss decreased (4.522735 --> 4.440095). Saving model ...		
Epoch: 8	Training Loss: 4.105937	Validation Loss: 4.399250
Validation loss decreased (4.440095 --> 4.399250). Saving model ...		
Epoch: 9	Training Loss: 4.016523	Validation Loss: 4.176144
Validation loss decreased (4.399250 --> 4.176144). Saving model ...		
Epoch: 10	Training Loss: 3.911892	Validation Loss: 4.216344
Epoch: 11	Training Loss: 3.837239	Validation Loss: 4.367939
Epoch: 12	Training Loss: 3.736571	Validation Loss: 4.039827
Validation loss decreased (4.176144 --> 4.039827). Saving model ...		
Epoch: 13	Training Loss: 3.646937	Validation Loss: 3.958032
Validation loss decreased (4.039827 --> 3.958032). Saving model ...		
Epoch: 14	Training Loss: 3.554220	Validation Loss: 4.116770
Epoch: 15	Training Loss: 3.478735	Validation Loss: 3.948525
Validation loss decreased (3.958032 --> 3.948525). Saving model ...		
Epoch: 16	Training Loss: 3.363377	Validation Loss: 3.884706
Validation loss decreased (3.948525 --> 3.884706). Saving model ...		
Epoch: 17	Training Loss: 3.274325	Validation Loss: 3.838474
Validation loss decreased (3.884706 --> 3.838474). Saving model ...		
Epoch: 18	Training Loss: 3.169669	Validation Loss: 3.942212
Epoch: 19	Training Loss: 3.074814	Validation Loss: 4.224119
Epoch: 20	Training Loss: 2.954479	Validation Loss: 3.864537
Epoch: 21	Training Loss: 2.834234	Validation Loss: 3.882875
Epoch: 22	Training Loss: 2.738972	Validation Loss: 3.845746
Epoch: 23	Training Loss: 2.632764	Validation Loss: 3.886914
Epoch: 24	Training Loss: 2.496515	Validation Loss: 3.909535
Epoch: 25	Training Loss: 2.390710	Validation Loss: 3.821160
Validation loss decreased (3.838474 --> 3.821160). Saving model ...		
Epoch: 26	Training Loss: 2.253152	Validation Loss: 3.874713
Epoch: 27	Training Loss: 2.142426	Validation Loss: 3.911588
Epoch: 28	Training Loss: 2.034972	Validation Loss: 3.964688
Epoch: 29	Training Loss: 1.899691	Validation Loss: 3.840306
Epoch: 30	Training Loss: 1.748688	Validation Loss: 3.940062
Epoch: 31	Training Loss: 1.647949	Validation Loss: 3.963384
Epoch: 32	Training Loss: 1.519266	Validation Loss: 4.065378
Epoch: 33	Training Loss: 1.422767	Validation Loss: 4.038896
Epoch: 34	Training Loss: 1.363646	Validation Loss: 4.109205
Epoch: 35	Training Loss: 1.255184	Validation Loss: 4.021675
Epoch: 36	Training Loss: 1.122894	Validation Loss: 4.005271
Epoch: 37	Training Loss: 1.043092	Validation Loss: 4.046155
Epoch: 38	Training Loss: 0.984172	Validation Loss: 4.080964
Epoch: 39	Training Loss: 0.863840	Validation Loss: 4.107561

Epoch: 40	Training Loss: 0.809669	Validation Loss: 4.238978
Epoch: 41	Training Loss: 0.750227	Validation Loss: 4.213589
Epoch: 42	Training Loss: 0.676087	Validation Loss: 4.331332
Epoch: 43	Training Loss: 0.625549	Validation Loss: 4.256974
Epoch: 44	Training Loss: 0.573893	Validation Loss: 4.268788
Epoch: 45	Training Loss: 0.554618	Validation Loss: 4.332600
Epoch: 46	Training Loss: 0.515905	Validation Loss: 4.343984
Epoch: 47	Training Loss: 0.476054	Validation Loss: 4.472133
Epoch: 48	Training Loss: 0.450636	Validation Loss: 4.334009
Epoch: 49	Training Loss: 0.407509	Validation Loss: 4.401799
Epoch: 50	Training Loss: 0.392475	Validation Loss: 4.512152
Epoch: 51	Training Loss: 0.369875	Validation Loss: 4.464147
Epoch: 52	Training Loss: 0.351269	Validation Loss: 4.500285
Epoch: 53	Training Loss: 0.315536	Validation Loss: 4.584879
Epoch: 54	Training Loss: 0.319270	Validation Loss: 4.595180
Epoch: 55	Training Loss: 0.306266	Validation Loss: 4.578832
Epoch: 56	Training Loss: 0.282546	Validation Loss: 4.659887
Epoch: 57	Training Loss: 0.267772	Validation Loss: 4.656134
Epoch: 58	Training Loss: 0.232334	Validation Loss: 4.637737
Epoch: 59	Training Loss: 0.262685	Validation Loss: 4.702690
Epoch: 60	Training Loss: 0.223369	Validation Loss: 4.870451
Epoch: 61	Training Loss: 0.225377	Validation Loss: 4.781228
Epoch: 62	Training Loss: 0.215229	Validation Loss: 4.855132
Epoch: 63	Training Loss: 0.202986	Validation Loss: 4.796535
Epoch: 64	Training Loss: 0.207825	Validation Loss: 4.838460
Epoch: 65	Training Loss: 0.177890	Validation Loss: 4.919020
Epoch: 66	Training Loss: 0.180397	Validation Loss: 4.842749
Epoch: 67	Training Loss: 0.173632	Validation Loss: 4.844027
Epoch: 68	Training Loss: 0.178193	Validation Loss: 4.830736
Epoch: 69	Training Loss: 0.159565	Validation Loss: 4.857112
Epoch: 70	Training Loss: 0.160898	Validation Loss: 4.801183
Epoch: 71	Training Loss: 0.169295	Validation Loss: 5.008532
Epoch: 72	Training Loss: 0.150117	Validation Loss: 4.862096
Epoch: 73	Training Loss: 0.139683	Validation Loss: 4.878680
Epoch: 74	Training Loss: 0.152478	Validation Loss: 4.808571
Epoch: 75	Training Loss: 0.134955	Validation Loss: 4.871299
Epoch: 76	Training Loss: 0.125851	Validation Loss: 4.848272
Epoch: 77	Training Loss: 0.126269	Validation Loss: 4.970186
Epoch: 78	Training Loss: 0.128774	Validation Loss: 4.977873
Epoch: 79	Training Loss: 0.119772	Validation Loss: 4.924303
Epoch: 80	Training Loss: 0.107659	Validation Loss: 5.115578
Epoch: 81	Training Loss: 0.109833	Validation Loss: 5.072667
Epoch: 82	Training Loss: 0.109573	Validation Loss: 4.979229
Epoch: 83	Training Loss: 0.107458	Validation Loss: 5.177440
Epoch: 84	Training Loss: 0.106956	Validation Loss: 5.068801
Epoch: 85	Training Loss: 0.102681	Validation Loss: 4.993513
Epoch: 86	Training Loss: 0.106469	Validation Loss: 5.100525
Epoch: 87	Training Loss: 0.093274	Validation Loss: 4.992168

Epoch: 88	Training Loss: 0.094210	Validation Loss: 5.056087
Epoch: 89	Training Loss: 0.100516	Validation Loss: 5.055184
Epoch: 90	Training Loss: 0.103470	Validation Loss: 5.180474
Epoch: 91	Training Loss: 0.096631	Validation Loss: 5.026805
Epoch: 92	Training Loss: 0.082145	Validation Loss: 5.080624
Epoch: 93	Training Loss: 0.087390	Validation Loss: 5.103360
Epoch: 94	Training Loss: 0.084554	Validation Loss: 5.028766
Epoch: 95	Training Loss: 0.074126	Validation Loss: 5.183977
Epoch: 96	Training Loss: 0.088981	Validation Loss: 5.148665
Epoch: 97	Training Loss: 0.086385	Validation Loss: 5.187917
Epoch: 98	Training Loss: 0.076803	Validation Loss: 5.196313
Epoch: 99	Training Loss: 0.081358	Validation Loss: 5.070851
Epoch: 100	Training Loss: 0.085773	Validation Loss: 5.088829

Out[45]: <All keys matched successfully>

### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

In [37]: `def test(loaders, model, criterion, use_cuda):`

```
    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))
```

```
print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))
```

```
In [13]: # call test function
         test(loaders_scratch, model_scratch, criterion_scratch, cuda_available)
```

Test Loss: 3.817867

Test Accuracy: 13% (116/836)

---

## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

### 1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [46]: ## TODO: Specify data loaders
         loaders_transfer = dict()
         loaders_transfer['train'] = trainloader
         loaders_transfer['test'] = testloader
         loaders_transfer['valid'] = validloader
```

### 1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable model\_transfer.

```
In [25]: class classifier(nn.Module):
         def __init__(self, c_0_inputs):
             super().__init__()
             self.fc1 = nn.Linear(c_0_inputs, 1000)
             self.fc2 = nn.Linear(1000, 1000)
             self.fc3 = nn.Linear(1000, 500)
             self.fc4 = nn.Linear(500, 133)

             self.batch_norm2 = nn.BatchNorm1d(num_features=1000)
             self.batch_norm3 = nn.BatchNorm1d(num_features=500)
             self.dropout = nn.Dropout(0.25)
```

```

def forward(self, x):
    # make sure input tensor is flattened
    x = x.view(x.size(0), -1)

    x = self.fc1(x)
    x = F.relu(x)
    x = self.dropout(x)
    x = self.fc2(x)
    x = F.relu(x)
    x = self.dropout(x)
    x = self.batch_norm2(x)
    x = self.fc3(x)
    x = F.relu(x)
    x = self.batch_norm3(x)
    x = self.fc4(x)

    return x

```

```

In [26]: ## TODO: Specify model architecture
vgg16 = models.vgg16(pretrained=True)

# Freeze training for all "features" layers
for param in vgg16.features.parameters():
    param.requires_grad = False

# I will replace all the classification layers.

c_0_inputs = vgg16.classifier[0].in_features
vgg16.classifier = classifier(c_0_inputs)

print(vgg16.classifier)

if cuda_available:
    model_transfer = vgg16.cuda()

classifier(
    (fc1): Linear(in_features=25088, out_features=1000, bias=True)
    (fc2): Linear(in_features=1000, out_features=1000, bias=True)
    (fc3): Linear(in_features=1000, out_features=500, bias=True)
    (fc4): Linear(in_features=500, out_features=133, bias=True)
    (batch_norm2): BatchNorm1d(1000, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (batch_norm3): BatchNorm1d(500, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (dropout): Dropout(p=0.25, inplace=False)
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Reviewer comments 1: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

In this exercise, I kept vgg16 feature extraction phase and replaced the learning phase entirely by replacing it with the learning phase of the cnn that I created earlier. I did this, in order to test my hypothesis, that the feature extraction phase in the custom made cnn was not properly design. Features in the feature extraction phase were kept and in the learning phase were retrained. Results were very promising, test accuracy of 84% was achieved.

#### 1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [39]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(vgg16.classifier.parameters(), lr=0.001)
```

#### 1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [25]: # train the transfer model
         loss_batch, valid_loss, model_transfer = train(100, loaders_transfer, model_transfer,
         criterion_transfer, cuda_available, 'model_scratch3.pt')

         # load the model that got the best validation accuracy (uncomment the line below)
         model_transfer.load_state_dict(torch.load('model_scratch3.pt'))
```

```
Epoch: 1      Training Loss: 4.597066      Validation Loss: 4.071239
Validation loss decreased (inf --> 4.071239).  Saving model ...
Epoch: 2      Training Loss: 3.832786      Validation Loss: 3.294235
Validation loss decreased (4.071239 --> 3.294235).  Saving model ...
Epoch: 3      Training Loss: 3.225330      Validation Loss: 2.744397
Validation loss decreased (3.294235 --> 2.744397).  Saving model ...
Epoch: 4      Training Loss: 2.743811      Validation Loss: 2.333085
Validation loss decreased (2.744397 --> 2.333085).  Saving model ...
Epoch: 5      Training Loss: 2.377515      Validation Loss: 2.024208
Validation loss decreased (2.333085 --> 2.024208).  Saving model ...
Epoch: 6      Training Loss: 2.087905      Validation Loss: 1.792346
Validation loss decreased (2.024208 --> 1.792346).  Saving model ...
Epoch: 7      Training Loss: 1.852640      Validation Loss: 1.595844
Validation loss decreased (1.792346 --> 1.595844).  Saving model ...
Epoch: 8      Training Loss: 1.657355      Validation Loss: 1.470441
Validation loss decreased (1.595844 --> 1.470441).  Saving model ...
Epoch: 9      Training Loss: 1.489828      Validation Loss: 1.308649
Validation loss decreased (1.470441 --> 1.308649).  Saving model ...
Epoch: 10     Training Loss: 1.368397      Validation Loss: 1.211461
```

Validation loss decreased (1.308649 --> 1.211461). Saving model ...

Epoch: 11	Training Loss: 1.240888	Validation Loss: 1.113330
-----------	-------------------------	---------------------------

Validation loss decreased (1.211461 --> 1.113330). Saving model ...

Epoch: 12	Training Loss: 1.128628	Validation Loss: 1.017311
-----------	-------------------------	---------------------------

Validation loss decreased (1.113330 --> 1.017311). Saving model ...

Epoch: 13	Training Loss: 1.043189	Validation Loss: 0.965731
-----------	-------------------------	---------------------------

Validation loss decreased (1.017311 --> 0.965731). Saving model ...

Epoch: 14	Training Loss: 0.956369	Validation Loss: 0.918134
-----------	-------------------------	---------------------------

Validation loss decreased (0.965731 --> 0.918134). Saving model ...

Epoch: 15	Training Loss: 0.885994	Validation Loss: 0.872082
-----------	-------------------------	---------------------------

Validation loss decreased (0.918134 --> 0.872082). Saving model ...

Epoch: 16	Training Loss: 0.817838	Validation Loss: 0.828464
-----------	-------------------------	---------------------------

Validation loss decreased (0.872082 --> 0.828464). Saving model ...

Epoch: 17	Training Loss: 0.768766	Validation Loss: 0.795059
-----------	-------------------------	---------------------------

Validation loss decreased (0.828464 --> 0.795059). Saving model ...

Epoch: 18	Training Loss: 0.702580	Validation Loss: 0.771254
-----------	-------------------------	---------------------------

Validation loss decreased (0.795059 --> 0.771254). Saving model ...

Epoch: 19	Training Loss: 0.663819	Validation Loss: 0.726378
-----------	-------------------------	---------------------------

Validation loss decreased (0.771254 --> 0.726378). Saving model ...

Epoch: 20	Training Loss: 0.616127	Validation Loss: 0.679994
-----------	-------------------------	---------------------------

Validation loss decreased (0.726378 --> 0.679994). Saving model ...

Epoch: 21	Training Loss: 0.565106	Validation Loss: 0.686237
-----------	-------------------------	---------------------------

Epoch: 22	Training Loss: 0.533259	Validation Loss: 0.677633
-----------	-------------------------	---------------------------

Validation loss decreased (0.679994 --> 0.677633). Saving model ...

Epoch: 23	Training Loss: 0.499670	Validation Loss: 0.669004
-----------	-------------------------	---------------------------

Validation loss decreased (0.677633 --> 0.669004). Saving model ...

Epoch: 24	Training Loss: 0.472985	Validation Loss: 0.642696
-----------	-------------------------	---------------------------

Validation loss decreased (0.669004 --> 0.642696). Saving model ...

Epoch: 25	Training Loss: 0.442701	Validation Loss: 0.632628
-----------	-------------------------	---------------------------

Validation loss decreased (0.642696 --> 0.632628). Saving model ...

Epoch: 26	Training Loss: 0.418468	Validation Loss: 0.604674
-----------	-------------------------	---------------------------

Validation loss decreased (0.632628 --> 0.604674). Saving model ...

Epoch: 27	Training Loss: 0.389862	Validation Loss: 0.598525
-----------	-------------------------	---------------------------

Validation loss decreased (0.604674 --> 0.598525). Saving model ...

Epoch: 28	Training Loss: 0.362864	Validation Loss: 0.582931
-----------	-------------------------	---------------------------

Validation loss decreased (0.598525 --> 0.582931). Saving model ...

Epoch: 29	Training Loss: 0.341189	Validation Loss: 0.589163
-----------	-------------------------	---------------------------

Epoch: 30	Training Loss: 0.330714	Validation Loss: 0.564195
-----------	-------------------------	---------------------------

Validation loss decreased (0.582931 --> 0.564195). Saving model ...

Epoch: 31	Training Loss: 0.310584	Validation Loss: 0.572365
-----------	-------------------------	---------------------------

Epoch: 32	Training Loss: 0.292987	Validation Loss: 0.559096
-----------	-------------------------	---------------------------

Validation loss decreased (0.564195 --> 0.559096). Saving model ...

Epoch: 33	Training Loss: 0.274128	Validation Loss: 0.556561
-----------	-------------------------	---------------------------

Validation loss decreased (0.559096 --> 0.556561). Saving model ...

Epoch: 34	Training Loss: 0.263129	Validation Loss: 0.539932
-----------	-------------------------	---------------------------

Validation loss decreased (0.556561 --> 0.539932). Saving model ...

Epoch: 35	Training Loss: 0.246184	Validation Loss: 0.541683
-----------	-------------------------	---------------------------

Epoch: 36	Training Loss: 0.233690	Validation Loss: 0.531535
-----------	-------------------------	---------------------------

Validation loss decreased (0.539932 --> 0.531535). Saving model ...

Epoch: 37	Training Loss: 0.224427	Validation Loss: 0.531884
Epoch: 38	Training Loss: 0.213584	Validation Loss: 0.533269
Epoch: 39	Training Loss: 0.205544	Validation Loss: 0.526437
Validation loss decreased (0.531535 --> 0.526437). Saving model ...		
Epoch: 40	Training Loss: 0.189440	Validation Loss: 0.515294
Validation loss decreased (0.526437 --> 0.515294). Saving model ...		
Epoch: 41	Training Loss: 0.180155	Validation Loss: 0.523311
Epoch: 42	Training Loss: 0.176149	Validation Loss: 0.510906
Validation loss decreased (0.515294 --> 0.510906). Saving model ...		
Epoch: 43	Training Loss: 0.165825	Validation Loss: 0.512147
Epoch: 44	Training Loss: 0.161219	Validation Loss: 0.508459
Validation loss decreased (0.510906 --> 0.508459). Saving model ...		
Epoch: 45	Training Loss: 0.152159	Validation Loss: 0.511513
Epoch: 46	Training Loss: 0.145803	Validation Loss: 0.512393
Epoch: 47	Training Loss: 0.140234	Validation Loss: 0.501484
Validation loss decreased (0.508459 --> 0.501484). Saving model ...		
Epoch: 48	Training Loss: 0.130443	Validation Loss: 0.495359
Validation loss decreased (0.501484 --> 0.495359). Saving model ...		
Epoch: 49	Training Loss: 0.129843	Validation Loss: 0.492279
Validation loss decreased (0.495359 --> 0.492279). Saving model ...		
Epoch: 50	Training Loss: 0.126662	Validation Loss: 0.486382
Validation loss decreased (0.492279 --> 0.486382). Saving model ...		
Epoch: 51	Training Loss: 0.121543	Validation Loss: 0.487514
Epoch: 52	Training Loss: 0.115404	Validation Loss: 0.473096
Validation loss decreased (0.486382 --> 0.473096). Saving model ...		
Epoch: 53	Training Loss: 0.109280	Validation Loss: 0.480492
Epoch: 54	Training Loss: 0.108048	Validation Loss: 0.495054
Epoch: 55	Training Loss: 0.100715	Validation Loss: 0.488889
Epoch: 56	Training Loss: 0.099200	Validation Loss: 0.493808
Epoch: 57	Training Loss: 0.096568	Validation Loss: 0.484988
Epoch: 58	Training Loss: 0.091654	Validation Loss: 0.491377
Epoch: 59	Training Loss: 0.090946	Validation Loss: 0.502080
Epoch: 60	Training Loss: 0.087332	Validation Loss: 0.498340
Epoch: 61	Training Loss: 0.086527	Validation Loss: 0.494710
Epoch: 62	Training Loss: 0.081584	Validation Loss: 0.482120
Epoch: 63	Training Loss: 0.076770	Validation Loss: 0.489178
Epoch: 64	Training Loss: 0.075926	Validation Loss: 0.474426
Epoch: 65	Training Loss: 0.075897	Validation Loss: 0.495338
Epoch: 66	Training Loss: 0.071602	Validation Loss: 0.502745
Epoch: 67	Training Loss: 0.072047	Validation Loss: 0.480503
Epoch: 68	Training Loss: 0.068900	Validation Loss: 0.489096
Epoch: 69	Training Loss: 0.066317	Validation Loss: 0.480594
Epoch: 70	Training Loss: 0.066531	Validation Loss: 0.486492
Epoch: 71	Training Loss: 0.062390	Validation Loss: 0.476981
Epoch: 72	Training Loss: 0.059585	Validation Loss: 0.484700
Epoch: 73	Training Loss: 0.061383	Validation Loss: 0.472365
Validation loss decreased (0.473096 --> 0.472365). Saving model ...		



Epoch: 74	Training Loss: 0.058774	Validation Loss: 0.493007
Epoch: 75	Training Loss: 0.059019	Validation Loss: 0.482265
Epoch: 76	Training Loss: 0.056204	Validation Loss: 0.487423
Epoch: 77	Training Loss: 0.057195	Validation Loss: 0.479320
Epoch: 78	Training Loss: 0.054174	Validation Loss: 0.490018
Epoch: 79	Training Loss: 0.052546	Validation Loss: 0.481440
Epoch: 80	Training Loss: 0.049684	Validation Loss: 0.463954
Validation loss decreased (0.472365 --> 0.463954). Saving model ...		
Epoch: 81	Training Loss: 0.051010	Validation Loss: 0.483553
Epoch: 82	Training Loss: 0.048822	Validation Loss: 0.460831
Validation loss decreased (0.463954 --> 0.460831). Saving model ...		
Epoch: 83	Training Loss: 0.048803	Validation Loss: 0.477836
Epoch: 84	Training Loss: 0.046561	Validation Loss: 0.487093
Epoch: 85	Training Loss: 0.045667	Validation Loss: 0.476361
Epoch: 86	Training Loss: 0.046498	Validation Loss: 0.478918
Epoch: 87	Training Loss: 0.044026	Validation Loss: 0.470675
Epoch: 88	Training Loss: 0.045211	Validation Loss: 0.464756
Epoch: 89	Training Loss: 0.044179	Validation Loss: 0.475361
Epoch: 90	Training Loss: 0.042325	Validation Loss: 0.483828
Epoch: 91	Training Loss: 0.042452	Validation Loss: 0.463791
Epoch: 92	Training Loss: 0.040530	Validation Loss: 0.461863
Epoch: 93	Training Loss: 0.041583	Validation Loss: 0.473152
Epoch: 94	Training Loss: 0.039412	Validation Loss: 0.461955
Epoch: 95	Training Loss: 0.039283	Validation Loss: 0.476914
Epoch: 96	Training Loss: 0.038170	Validation Loss: 0.462210
Epoch: 97	Training Loss: 0.038758	Validation Loss: 0.484857
Epoch: 98	Training Loss: 0.037592	Validation Loss: 0.476753
Epoch: 99	Training Loss: 0.037244	Validation Loss: 0.469756
Epoch: 100	Training Loss: 0.035534	Validation Loss: 0.473606

Out[25]: <All keys matched successfully>

### (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

In [21]: `model_transfer.load_state_dict(torch.load('model_scratch3.pt'))`

Out[21]: <All keys matched successfully>

In [13]: `test(loaders_transfer, model_transfer, criterion_transfer, cuda_available)`

Test Loss: 0.501152

Test Accuracy: 84% (706/836)

My comments 2: After modifying the loaders, I directly applied the transfer-learned model (model\_scratch3.pt), which previously achieved a precision of 84% to the new test set data loader. The new achieved accuracy is of 49%. The code for this the new test is show in the following cell.

```
In [38]: model_transfer.load_state_dict(torch.load('model_scratch3.pt'))
         test(loaders_transfer, model_transfer, criterion_transfer, cuda_available)
```

Test Loss: 2.046817

Test Accuracy: 49% (411/836)

My comments 2: Results degradate by a factor of 1.7!. It seems that image pre-processing phases corresponding to random horizontal flip, rotation, and center cropping significantly improved test accuracy. As described earlier, model scratch 3 was not retrained using a validation loader that omitted the previously mention transformations. To analyze if the model quality improves, it is retrained and retested. See code in the following cells.

```
In [47]: # train the transfer model
         loss_batch, valid_loss, model_transfer = train(100, loaders_transfer, model_transfer,
         criterion_transfer, cuda_available, 'model_scratch4.pt')
```

```
Epoch: 1      Training Loss: 0.047734      Validation Loss: 0.469103
Validation loss decreased (inf --> 0.469103). Saving model ...
Epoch: 2      Training Loss: 0.048244      Validation Loss: 0.463061
Validation loss decreased (0.469103 --> 0.463061). Saving model ...
Epoch: 3      Training Loss: 0.047332      Validation Loss: 0.467075
Epoch: 4      Training Loss: 0.043777      Validation Loss: 0.466281
Epoch: 5      Training Loss: 0.044889      Validation Loss: 0.468572
Epoch: 6      Training Loss: 0.045816      Validation Loss: 0.466513
Epoch: 7      Training Loss: 0.042952      Validation Loss: 0.474006
Epoch: 8      Training Loss: 0.042780      Validation Loss: 0.473524
Epoch: 9      Training Loss: 0.041193      Validation Loss: 0.463608
Epoch: 10     Training Loss: 0.040998      Validation Loss: 0.462575
Validation loss decreased (0.463061 --> 0.462575). Saving model ...
Epoch: 11     Training Loss: 0.041538      Validation Loss: 0.471334
Epoch: 12     Training Loss: 0.039427      Validation Loss: 0.473825
Epoch: 13     Training Loss: 0.040018      Validation Loss: 0.468763
Epoch: 14     Training Loss: 0.039896      Validation Loss: 0.477060
Epoch: 15     Training Loss: 0.037594      Validation Loss: 0.470795
Epoch: 16     Training Loss: 0.036319      Validation Loss: 0.463526
Epoch: 17     Training Loss: 0.035309      Validation Loss: 0.461256
Validation loss decreased (0.462575 --> 0.461256). Saving model ...
Epoch: 18     Training Loss: 0.034659      Validation Loss: 0.466602
Epoch: 19     Training Loss: 0.035590      Validation Loss: 0.464091
Epoch: 20     Training Loss: 0.033860      Validation Loss: 0.467963
Epoch: 21     Training Loss: 0.035094      Validation Loss: 0.465130
Epoch: 22     Training Loss: 0.032585      Validation Loss: 0.470218
```

Epoch: 23	Training Loss: 0.033375	Validation Loss: 0.464225
Epoch: 24	Training Loss: 0.032623	Validation Loss: 0.460103
Validation loss decreased (0.461256 --> 0.460103). Saving model ...		
Epoch: 25	Training Loss: 0.031729	Validation Loss: 0.460704
Epoch: 26	Training Loss: 0.032139	Validation Loss: 0.466804
Epoch: 27	Training Loss: 0.030654	Validation Loss: 0.461972
Epoch: 28	Training Loss: 0.032488	Validation Loss: 0.461988
Epoch: 29	Training Loss: 0.030130	Validation Loss: 0.467273
Epoch: 30	Training Loss: 0.031967	Validation Loss: 0.461604
Epoch: 31	Training Loss: 0.029690	Validation Loss: 0.464949
Epoch: 32	Training Loss: 0.029583	Validation Loss: 0.465034
Epoch: 33	Training Loss: 0.029677	Validation Loss: 0.471237
Epoch: 34	Training Loss: 0.029178	Validation Loss: 0.470556
Epoch: 35	Training Loss: 0.029086	Validation Loss: 0.473892
Epoch: 36	Training Loss: 0.028025	Validation Loss: 0.470112
Epoch: 37	Training Loss: 0.028330	Validation Loss: 0.471677
Epoch: 38	Training Loss: 0.026484	Validation Loss: 0.468820
Epoch: 39	Training Loss: 0.026376	Validation Loss: 0.468450
Epoch: 40	Training Loss: 0.027421	Validation Loss: 0.473695
Epoch: 41	Training Loss: 0.025220	Validation Loss: 0.473905
Epoch: 42	Training Loss: 0.026333	Validation Loss: 0.468383
Epoch: 43	Training Loss: 0.027277	Validation Loss: 0.473397
Epoch: 44	Training Loss: 0.026784	Validation Loss: 0.470728
Epoch: 45	Training Loss: 0.023879	Validation Loss: 0.466043
Epoch: 46	Training Loss: 0.025111	Validation Loss: 0.472101
Epoch: 47	Training Loss: 0.025083	Validation Loss: 0.472508
Epoch: 48	Training Loss: 0.024165	Validation Loss: 0.465119
Epoch: 49	Training Loss: 0.024136	Validation Loss: 0.473288
Epoch: 50	Training Loss: 0.024657	Validation Loss: 0.474809
Epoch: 51	Training Loss: 0.023883	Validation Loss: 0.472117
Epoch: 52	Training Loss: 0.023963	Validation Loss: 0.476374
Epoch: 53	Training Loss: 0.023631	Validation Loss: 0.471976
Epoch: 54	Training Loss: 0.023861	Validation Loss: 0.474686
Epoch: 55	Training Loss: 0.022959	Validation Loss: 0.473814
Epoch: 56	Training Loss: 0.022279	Validation Loss: 0.478200
Epoch: 57	Training Loss: 0.021797	Validation Loss: 0.468825
Epoch: 58	Training Loss: 0.022152	Validation Loss: 0.478918
Epoch: 59	Training Loss: 0.022037	Validation Loss: 0.470872
Epoch: 60	Training Loss: 0.021203	Validation Loss: 0.473166
Epoch: 61	Training Loss: 0.021411	Validation Loss: 0.475494
Epoch: 62	Training Loss: 0.020483	Validation Loss: 0.470029
Epoch: 63	Training Loss: 0.021653	Validation Loss: 0.478640
Epoch: 64	Training Loss: 0.021306	Validation Loss: 0.474579
Epoch: 65	Training Loss: 0.021165	Validation Loss: 0.472056
Epoch: 66	Training Loss: 0.020645	Validation Loss: 0.467473
Epoch: 67	Training Loss: 0.020256	Validation Loss: 0.471857
Epoch: 68	Training Loss: 0.020742	Validation Loss: 0.468695
Epoch: 69	Training Loss: 0.020832	Validation Loss: 0.473682

Epoch: 70	Training Loss: 0.020289	Validation Loss: 0.468472
Epoch: 71	Training Loss: 0.019761	Validation Loss: 0.467030
Epoch: 72	Training Loss: 0.019826	Validation Loss: 0.471278
Epoch: 73	Training Loss: 0.019536	Validation Loss: 0.465755
Epoch: 74	Training Loss: 0.018913	Validation Loss: 0.472406
Epoch: 75	Training Loss: 0.018993	Validation Loss: 0.462947
Epoch: 76	Training Loss: 0.018453	Validation Loss: 0.470613
Epoch: 77	Training Loss: 0.018279	Validation Loss: 0.466213
Epoch: 78	Training Loss: 0.018502	Validation Loss: 0.471229
Epoch: 79	Training Loss: 0.018746	Validation Loss: 0.471243
Epoch: 80	Training Loss: 0.019141	Validation Loss: 0.467318
Epoch: 81	Training Loss: 0.018655	Validation Loss: 0.469253
Epoch: 82	Training Loss: 0.018614	Validation Loss: 0.472540
Epoch: 83	Training Loss: 0.017051	Validation Loss: 0.475628
Epoch: 84	Training Loss: 0.017579	Validation Loss: 0.472243
Epoch: 85	Training Loss: 0.019195	Validation Loss: 0.466429
Epoch: 86	Training Loss: 0.017705	Validation Loss: 0.470477
Epoch: 87	Training Loss: 0.017896	Validation Loss: 0.473380
Epoch: 88	Training Loss: 0.017020	Validation Loss: 0.480217
Epoch: 89	Training Loss: 0.017000	Validation Loss: 0.477352
Epoch: 90	Training Loss: 0.016935	Validation Loss: 0.477624
Epoch: 91	Training Loss: 0.017109	Validation Loss: 0.477153
Epoch: 92	Training Loss: 0.016254	Validation Loss: 0.472896
Epoch: 93	Training Loss: 0.017090	Validation Loss: 0.478364
Epoch: 94	Training Loss: 0.017183	Validation Loss: 0.480340
Epoch: 95	Training Loss: 0.018025	Validation Loss: 0.474726
Epoch: 96	Training Loss: 0.015349	Validation Loss: 0.473520
Epoch: 97	Training Loss: 0.016989	Validation Loss: 0.477230
Epoch: 98	Training Loss: 0.016524	Validation Loss: 0.474687
Epoch: 99	Training Loss: 0.016236	Validation Loss: 0.476214
Epoch: 100	Training Loss: 0.015955	Validation Loss: 0.467174

```
In [48]: model_transfer.load_state_dict(torch.load('model_scratch4.pt'))
        test(loaders_transfer, model_transfer, criterion_transfer, cuda_available)
```

Test Loss: 0.566806

Test Accuracy: 84% (708/836)

My comments 2: After retraining the model, test accuracy was restore to 84%. It actually, improved slightly from a ratio of 706/836 to 708/836.

### 1.1.16 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```

In [213]: # HELPER function. Gets the workload labels (unused, dont recomend this function)
def get_labels_from_loader(loader, name):
    items = [items for (_, items) in enumerate(loader[name])]
    lists = [item[1].tolist() for item in items]
    usval = sorted(set(sum(lists, []))) # unique sorted values indexed from 0 onward
    labels = [int(l) for l in usval]
    return labels

loader_class_labels = get_labels_from_loader(loaders_transfer, 'train')

In [22]: def get_labels_from_filesystem(img_working_dir):
    dictionary = dict()
    files = os.listdir(img_working_dir)
    for file in files:
        key, val = file.split('.')
        dictionary[int(key)] = val
    return dictionary

class_labels = get_labels_from_filesystem(os.getcwd() + "\\dogImages\\test")

In [24]: # HELPER function: given a file path, display the image
def display_img(img_path, name):
    params = {"text.color" : "blue"}
    img = Image.open(img_path)
    plt.imshow(img)
    plt.title(name, pad=10)
    plt.show()
    plt.rcParams.update(params)

In [324]: def show_predictions(model, loader, train_on_gpu, labels):
    # obtain one batch of test images
    dataiter = iter(loader)
    images, _ = dataiter.next() # at a later time analyze labels
    images.numpy()
    imgs = images

    # move model inputs to cuda, if GPU available
    #if train_on_gpu:
    #    images = images.cuda()
    if cuda_available:
        images = images.cuda()
        model = model.cuda()

    # get sample outputs
    output = model(images)

    # convert output probabilities to predicted class

```

```

_, preds_tensor = torch.max(output, 1)

if train_on_gpu:
    preds = np.squeeze(preds_tensor.cpu().numpy())
else:
    preds = np.squeeze(preds_tensor.numpy())

# plot the images in the batch, along with predicted and true labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
    img = imgs[idx]
    img = (img-torch.min(img))/(torch.max(img)-torch.min(img)) #apply min-max no
    plt.imshow(np.transpose(img, (1, 2, 0)))
    k = preds[idx]+1
    ax.set_title("{}".format(labels[k]))
    #ax.set_title("{} ({}).format(classes[preds[idx]], classes[labels[idx]]),
    # color=("green" if preds[idx]==labels[idx].item() else "red"))

```

In [20]: *### TODO: Write a function that takes a path to an image as input  
 ### and returns the dog breed that is predicted by the model.*

```

# list of class names by index, i.e. a name can be accessed like class_names[0]
def Dog_breed_detector(img_path):
    '''
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    '''
    # load and prepare the model
    model = models.vgg16(pretrained=True)

    # Freeze training for all "features" layers
    for param in model.features.parameters():
        param.requires_grad = False

    c_0_inputs = model.classifier[0].in_features
    model.classifier = classifier(c_0_inputs)

    model.load_state_dict(torch.load('model_scratch3.pt'))
    if cuda_available:
        model = model.cuda()
    model.eval()

```

```

img = Image.open(img_path)

# define image transformations
transformations = transforms.Compose([
    transforms.Resize(size=224),
    transforms.CenterCrop((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]))

# apply the image transformations
img = transformations(img).unsqueeze_(0)

# map to gpu
if cuda_available:
    img = img.cuda()

output = model(img)

_, pred = torch.max(output, 1)

pred = np.squeeze(pred.numpy()) if not cuda_available else np.squeeze(pred.cpu()).item()
return int(pred) + 1

In [25]: # Test code for transfer_predict. I did not use the class_labels, but the inverse tab
img_path = dog_files[2000]
cid = Dog_breed_detector(img_path) # previously named transfer_predict
labels = get_labels_from_filesystem(os.getcwd()+ "\\dogImages\\test")
name = "Predicted: label " + str(cid) + ", human readable name: " + labels[cid]
display_img(img_path, name)

```